

Heuristic Search

1

Heuristic Search

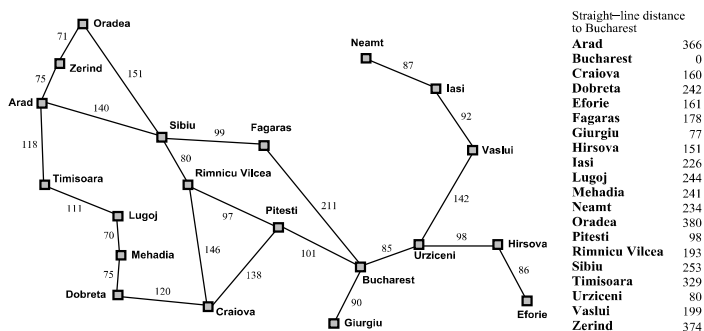
- **Heuristic or informed search** exploits additional knowledge about the problem that helps direct search to more promising paths.
- A **heuristic function**, $h(n)$, provides an estimate of the cost of the path from a given node to the closest goal state. Must be zero if node represents a goal state.
 - Example: Straight-line distance from current location to the goal location in a road navigation problem.
- Many search problems are NP-complete so in the worst case still have exponential time complexity; however a good heuristic can:
 - Find a solution for an average problem efficiently.
 - Find a reasonably good but not optimal solution efficiently.

2

Best-First Search

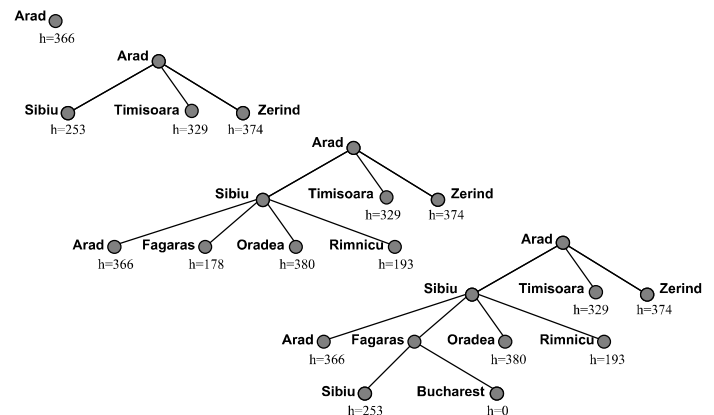
- At each step, best-first search sorts the queue according to a heuristic function.

function BEST-FIRST-SEARCH(*problem*, EVAL-FN) **returns** a solution sequence
inputs: *problem*, a problem
 EVAL-FN, an evaluation function
Queueing-FN ← a function that orders nodes by EVAL-FN
return GENERAL-SEARCH(*problem*, *Queueing-FN*)



3

Best-First Example



- Does not find shortest path to goal (through Rimnicu) since it is only focused on the cost remaining rather than the total cost.

4

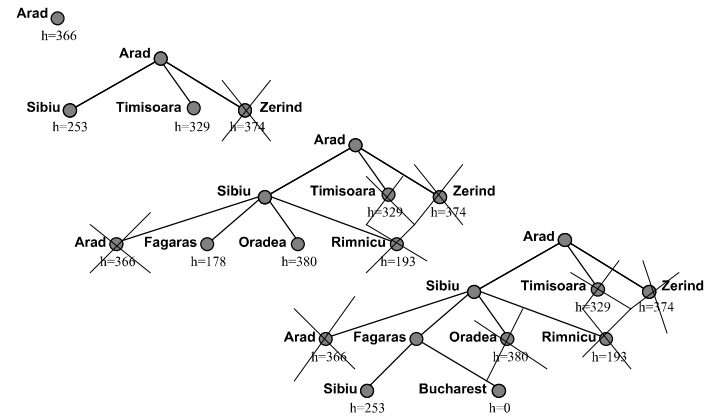
Best-First Properties

- Not complete, may follow infinite path if heuristic rates each state on such a path as the best option. Most reasonable heuristics will not cause this problem however.
- Worst case time complexity is still $O(b^m)$ where m is the maximum depth.
- Since must maintain a queue of all unexpanded states, space-complexity is also $O(b^m)$.
- However, a good heuristic will avoid this worst-case behavior for most problems.

5

Beam Search


- Space and time complexity of storing and sorting the complete queue can be too inefficient.
- Beam search trims queue to the best n options (n is called the **beam width**) at each point.
- Focuses search more but may eliminate solution even for finite search graphs
- Example for $n=2$.



6

Hill-Climbing

- Beam search with a beam-width of 1 is called **hill-climbing**.
- Pursues locally best option at each point, i.e. the best successor to the current best node.
- Subject to local maxima, plateaux, and ridges.



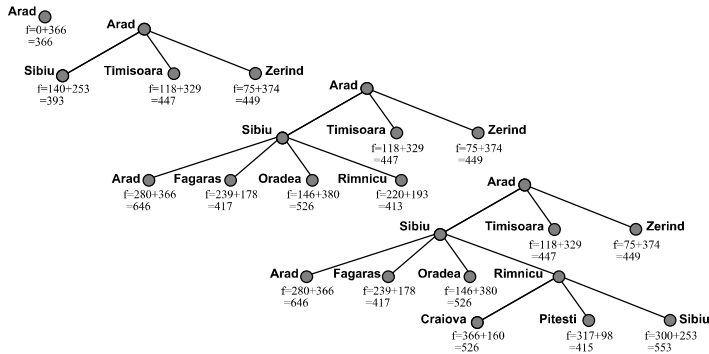
7

Minimizing Total Path Cost: A* Search

- A* combines features of uniform cost search (complete, optimal, inefficient) with best-first (incomplete, non-optimal, efficient).
- Sort queue by estimated total cost of the completion of a path.
$$f(n) = g(n) + h(n)$$
- If the heuristic function always underestimates the distance to the goal, it is said to be **admissible**.
- If h is admissible, then $f(n)$ never overestimates the actual cost of the best solution through n .

8

A* Example



- Finds the optimal path to Bucharest through Rimnicu and Pitesti

9

Optimality of A*

- If h is admissible, A* will always find a least cost path to the goal.

- Proof by contradiction:

Let G be an optimal goal state with a path cost f^*

Let G_2 be a suboptimal goal state supposedly found by A*

Let n be a current leaf node on an optimal path to G

Since h is admissible:

$$f^* \geq f(n)$$

If G_2 is chosen for expansion over n then:

$$f(n) \geq f(G_2)$$

Therefore:

$$f^* \geq f(G_2)$$

Since G_2 is a goal state, $h(G_2)=0$, therefore

$$f(G_2) = g(G_2)$$

$$f^* \geq g(G_2)$$

Therefore G_2 is optimal.

Contradiction.

Q.E.D.

10


Other Properties of A*

- A* is complete as long as
 - Branching factor is always finite
 - Every operator adds cost at least $\delta > 0$
- Time and space complexity still $O(b^m)$ in the worst case since must maintain and sort complete queue of unexplored options.
- However, with a good heuristic can find optimal solutions for many problems in reasonable time.
- Again, space complexity is a worse problem than time.


11

Heuristic Functions

- 8-puzzle search space
 - Typical solution length: 20 steps
 - Average branching factor: 3
 - Exhaustive search: $3^{20} = 3.5 \times 10^9$
 - Bound on unique states: $9! = 362,880$



Start State



Goal State

- Admissible Heuristics:
 - Number of tiles out of place (h_1): 7
 - City-block (Manhattan) distance (h_2): $2+3+3+2+4+2+0+2=18$

12

Heuristic Performance

- Experiments on sample problems can determine the number of nodes searched and CPU time for different strategies.

- One other useful measure is **effective branching factor**: If a method expands N nodes to find solution of depth d , and a uniform tree of depth d would require a branching factor of b^* to contain N nodes, the effective branching factor is b^*

$$N = 1 + b^* + (b^*)^2 + \dots + (b^*)^d$$

- Experimental Results on 8-puzzle problems

| d | Search Cost | | | Effective Branching Factor | | |
|-----|-------------|------------|------------|----------------------------|------------|------------|
| | IDS | $A^*(h_1)$ | $A^*(h_2)$ | IDS | $A^*(h_1)$ | $A^*(h_2)$ |
| 2 | 10 | 6 | 6 | 2.45 | 1.79 | 1.79 |
| 4 | 112 | 13 | 12 | 2.87 | 1.48 | 1.45 |
| 6 | 680 | 20 | 18 | 2.73 | 1.34 | 1.30 |
| 8 | 6384 | 39 | 25 | 2.80 | 1.33 | 1.24 |
| 10 | 47127 | 93 | 39 | 2.79 | 1.38 | 1.22 |
| 12 | 364404 | 227 | 73 | 2.78 | 1.42 | 1.24 |
| 14 | 3473941 | 539 | 113 | 2.83 | 1.44 | 1.23 |
| 16 | — | 1301 | 211 | — | 1.45 | 1.25 |
| 18 | — | 3056 | 363 | — | 1.46 | 1.26 |
| 20 | — | 7276 | 676 | — | 1.47 | 1.27 |
| 22 | — | 18094 | 1219 | — | 1.48 | 1.28 |
| 24 | — | 39135 | 1641 | — | 1.48 | 1.26 |

13

Quality of Heuristics

- Since A^* expands all nodes whose f value is less than that of an optimal solution, it is always better to use a heuristic with a higher value as long as it does not over-estimate.

- Therefore h_2 is uniformly better than h_1 , or h_2 **dominates** h_1 .

- A heuristic should also be easy to compute, otherwise the overhead of computing the heuristic could outweigh the time saved by reducing search (e.g. using full breadth-first search to estimate distance wouldn't help).

14

Inventing Heuristics

- Many good heuristics can be invented by considering relaxed versions of the problem (abstractions).

- For 8-puzzle:

A tile can move from square A to B if A is adjacent to B and B is blank

- A tile can move from square A to B if A is adjacent to B.
- A tile can move from square A to B if B is blank.
- A tile can move from square A to B.

- If there are a number of features that indicate a promising or unpromising state, a weighted sum of these features can be useful. Learning methods can be used to set weights.

15